Statistical philosophy and the language of significance

Jonathan Dushoff, McMaster University

U. Chicago, Oct 2018

P values

- We compare health indicators of children treated or not treated with vitamin A supplements

P values

- We compare health indicators of children treated or not treated with vitamin A supplements
- What does it mean if I find a “significant P value” for some effect in this experiment?

P values

- We compare health indicators of children treated or not treated with vitamin A supplements
- What does it mean if I find a “significant P value” for some effect in this experiment?
    - *

P values

- We compare health indicators of children treated or not treated with vitamin A supplements
- What does it mean if I find a “significant P value” for some effect in this experiment?
  - * The difference is unlikely to be due to chance

P values

- We compare health indicators of children treated or not treated with vitamin A supplements
- What does it mean if I find a “significant P value” for some effect in this experiment?
  - * The difference is unlikely to be due to chance
  - *

P values

- We compare health indicators of children treated or not treated with vitamin A supplements
- What does it mean if I find a “significant P value” for some effect in this experiment?
  - * The difference is unlikely to be due to chance
  - * So what! I already know vitamin A has strong effects on metabolism

P values

- We compare health indicators of children treated or not treated with vitamin A supplements
- What does it mean if I find a “significant P value” for some effect in this experiment?
    - * The difference is unlikely to be due to chance
    - * So what! I already know vitamin A has strong effects on metabolism
- If I’m certain that the true answer isn’t exactly zero, why do I want the P value anyway?

P values

- We compare health indicators of children treated or not treated with vitamin A supplements
- What does it mean if I find a “significant P value” for some effect in this experiment?
  - * The difference is unlikely to be due to chance
  - * So what! I already know vitamin A has strong effects on metabolism
- If I’m certain that the true answer isn’t exactly zero, why do I want the P value anyway?

Confidence intervals

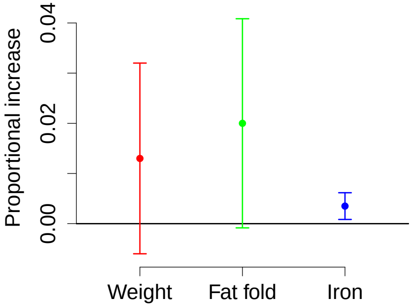

- What do these results mean?

Confidence intervals

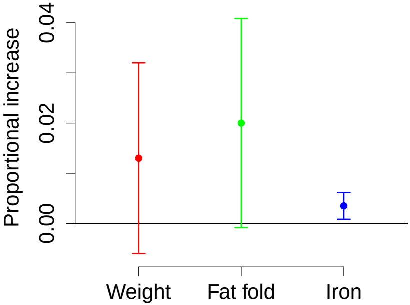

- What do these results mean?
- Which are significant?

Confidence intervals

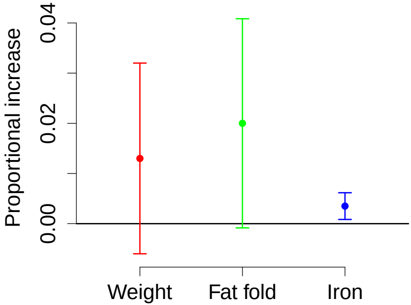

- What do these results mean?
- Which are significant?

Confidence intervals and P values

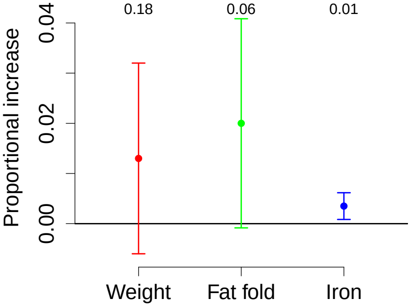

- A high P value means we can't see the *sign* of the effect clearly

Confidence intervals and P values

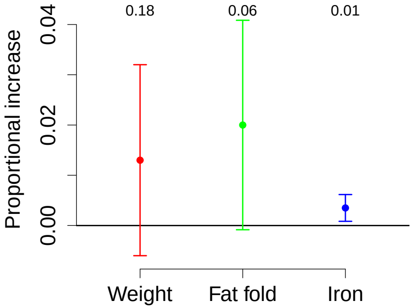

- A high P value means we can't see the *sign* of the effect clearly

- A low P value means we can

Confidence intervals and P values

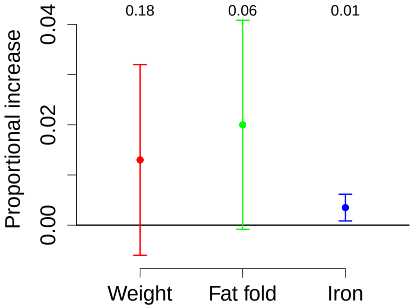

- A high P value means we can't see the *sign* of the effect clearly
- A low P value means we can

The meaning of P values

- More broadly, a P value measures whether we are seeing *something* clearly

The meaning of P values

- More broadly, a P value measures whether we are seeing *something* clearly
  - It's usually the sign (±) of some quantity, but doesn't need to be

The meaning of P values

- More broadly, a P value measures whether we are seeing *something* clearly
    - It's usually the sign (±) of some quantity, but doesn't need to be

Types of Error

- Type I (*False positive:*) concluding there is an effect when there isn't one

Types of Error

- Type I (*False positive:*) concluding there is an effect when there isn't one
    - This doesn't happen in biology. There is always an effect.

Types of Error

- Type I (*False positive:*) concluding there is an effect when there isn't one
  - This doesn't happen in biology. There is always an effect.
  - This is a defensible belief, and also an unfalsifiable one

Types of Error

- Type I (*False positive:*) concluding there is an effect when there isn't one
  - This doesn't happen in biology. There is always an effect.
  - This is a defensible belief, and also an unfalsifiable one
- Type II (*False negative:*) concluding there is no effect when there really is

Types of Error

- Type I (*False positive:*) concluding there is an effect when there isn't one
  - This doesn't happen in biology. There is always an effect.
  - This is a defensible belief, and also an unfalsifiable one
- Type II (*False negative:*) concluding there is no effect when there really is
  - This *should* never happen in biology, because we should never conclude there is no effect

Types of Error

- Type I (*False positive:*) concluding there is an effect when there isn't one
  - This doesn't happen in biology. There is always an effect.
  - This is a defensible belief, and also an unfalsifiable one
- Type II (*False negative:*) concluding there is no effect when there really is
  - This *should* never happen in biology, because we should never conclude there is no effect
  - In fact, it happens all the time

Types of Error

- Type I (*False positive:*) concluding there is an effect when there isn't one
  - This doesn't happen in biology. There is always an effect.
  - This is a defensible belief, and also an unfalsifiable one
- Type II (*False negative:*) concluding there is no effect when there really is
  - This *should* never happen in biology, because we should never conclude there is no effect
  - In fact, it happens all the time

Types of Error

- Type III Error is the error of using numerical codes for things that have perfectly good simple names

Types of Error

- Type III Error is the error of using numerical codes for things that have perfectly good simple names
- Just say "false positive" or "false negative" when possible

Types of Error

- Type III Error is the error of using numerical codes for things that have perfectly good simple names
- Just say “false positive” or “false negative” when possible

Errors in applied studies

- *Sign error:* if I think an effect is positive, when it's really negative (or vice versa)

Errors in applied studies

- *Sign error:* if I think an effect is positive, when it's really negative (or vice versa)
- *Magnitude error:* if I think an effect is small, when it's really large (or vice versa)

Errors in applied studies

- *Sign error:* if I think an effect is positive, when it's really negative (or vice versa)
- *Magnitude error:* if I think an effect is small, when it's really large (or vice versa)
- Confidence intervals clarify all of this

Errors in applied studies

- *Sign error:* if I think an effect is positive, when it's really negative (or vice versa)
- *Magnitude error:* if I think an effect is small, when it's really large (or vice versa)
- Confidence intervals clarify all of this

Errors in theoretical studies

- *False positive:* in the hypothetical case that the effect is exactly zero, what is the probability of falsely finding an effect?

Errors in theoretical studies

- *False positive:* in the hypothetical case that the effect is exactly zero, what is the probability of falsely finding an effect?
  - Should be less than or equal to my nominal significance value

Errors in theoretical studies

- *False positive:* in the hypothetical case that the effect is exactly zero, what is the probability of falsely finding an effect?
  - Should be less than or equal to my nominal significance value
  - This is the gold standard for statistical validity

Errors in theoretical studies

- *False positive:* in the hypothetical case that the effect is exactly zero, what is the probability of falsely finding an effect?
  - Should be less than or equal to my nominal significance value
  - This is the gold standard for statistical validity
- *False negative:* what is the probability of failing to find an effect that is there?

Errors in theoretical studies

- *False positive:* in the hypothetical case that the effect is exactly zero, what is the probability of falsely finding an effect?
  - Should be less than or equal to my nominal significance value
  - This is the gold standard for statistical validity
- *False negative:* what is the probability of failing to find an effect that is there?
  - Requires you specify a hypothetical effect *size*

Errors in theoretical studies

- *False positive:* in the hypothetical case that the effect is exactly zero, what is the probability of falsely finding an effect?
  - Should be less than or equal to my nominal significance value
  - This is the gold standard for statistical validity
- *False negative:* what is the probability of failing to find an effect that is there?
  - Requires you specify a hypothetical effect *size*
    - This is a *scientific* judgment

Errors in theoretical studies

- *False positive:* in the hypothetical case that the effect is exactly zero, what is the probability of falsely finding an effect?
  - Should be less than or equal to my nominal significance value
  - This is the gold standard for statistical validity
- *False negative:* what is the probability of failing to find an effect that is there?
  - Requires you specify a hypothetical effect *size*
    - This is a *scientific* judgment
  - This is a good way to analyze power

Errors in theoretical studies

- *False positive:* in the hypothetical case that the effect is exactly zero, what is the probability of falsely finding an effect?
  - Should be less than or equal to my nominal significance value
  - This is the gold standard for statistical validity
- *False negative:* what is the probability of failing to find an effect that is there?
  - Requires you specify a hypothetical effect *size*
    - This is a *scientific* judgment
  - This is a good way to analyze power
- You should do these analyses *before* you collect data, not after

Errors in theoretical studies

- *False positive:* in the hypothetical case that the effect is exactly zero, what is the probability of falsely finding an effect?
  - Should be less than or equal to my nominal significance value
  - This is the gold standard for statistical validity
- *False negative:* what is the probability of failing to find an effect that is there?
  - Requires you specify a hypothetical effect *size*
    - This is a *scientific* judgment
  - This is a good way to analyze power
- You should do these analyses *before* you collect data, not after

Low P values

- If I have a low P value I can see something clearly

Low P values

- If I have a low P value I can see something clearly
- But it's usually better to focus on what I see than the P value

Low P values

- If I have a low P value I can see something clearly
- But it's usually better to focus on what I see than the P value

High P values

- If I have a high P value, there is something I *don't* see clearly

High P values

- If I have a high P value, there is something I *don't* see clearly
- It *may be* because this effect is small

High P values

- If I have a high P value, there is something I *don't* see clearly
- It *may be* because this effect is small
- High P values should *not* be used to advance your conclusion

High P values

- If I have a high P value, there is something I *don't* see clearly
- It *may be* because this effect is small
- High P values should *not* be used to advance your conclusion

Are high P values evidence?

- What causes them?

Are high P values evidence?

- What causes them?
  - **Small differences**

Are high P values evidence?

- What causes them?
  - **Small differences**
  - Less data

Are high P values evidence?

- What causes them?
  - **Small differences**
  - Less data
  - More noise

Are high P values evidence?

- What causes them?
    - **Small differences**
    - Less data
    - More noise
    - An inappropriate model

Are high P values evidence?

- What causes them?
  - **Small differences**
  - Less data
  - More noise
  - An inappropriate model
- A lower P value means that your evidence for difference is better

Are high P values evidence?

- What causes them?
  - **Small differences**
  - Less data
  - More noise
  - An inappropriate model
- A lower P value means that your evidence for difference is better
- A higher P value means that your evidence for similarity is better – or worse!

Are high P values evidence?

- What causes them?
  - **Small differences**
  - Less data
  - More noise
  - An inappropriate model
- A lower P value means that your evidence for difference is better
- A higher P value means that your evidence for similarity is better – or worse!

Annualized flu deaths

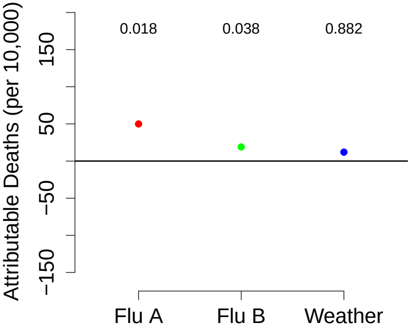

- Why is weather not causing deaths at this time scale?

Annualized flu deaths

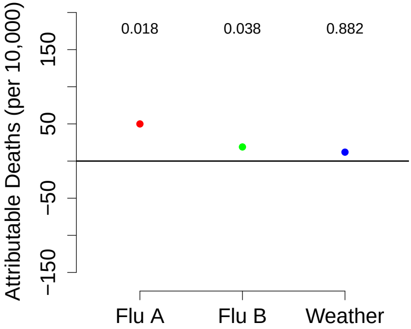

- Why is weather not causing deaths at this time scale?

… with confidence intervals

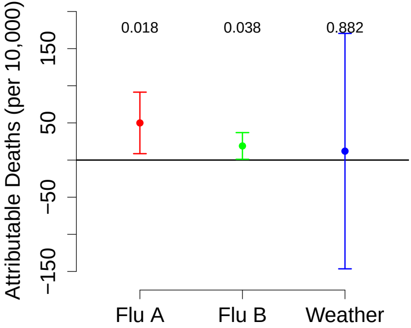

- **Never** say: A is significant and B isn't, so $A > B$

… with confidence intervals

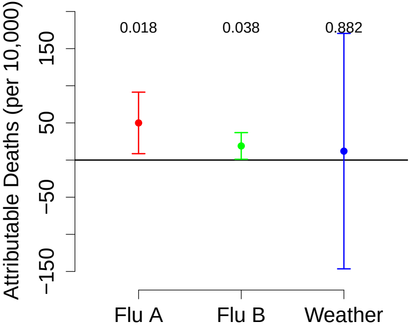

- **Never** say: A is significant and B isn't, so $A > B$
- **Instead:** Construct a statistic for the hypothesis $A > B$

... with confidence intervals

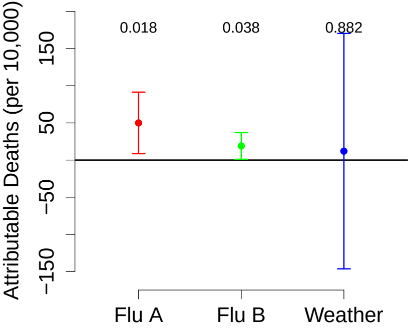

- **Never** say: A is significant and B isn't, so $A > B$
- **Instead:** Construct a statistic for the hypothesis $A > B$

Small effects

- Evidence that an effect is small:

Small effects

- Evidence that an effect is small:
  - We can’t even tell the sign!

Small effects

- Evidence that an effect is small:
    - We can't even tell the sign!
    - It looks small

Small effects

- Evidence that an effect is small:
    - We can’t even tell the sign!
    - It looks small
    - We are confident that it’s small

Small effects

- Evidence that an effect is small:
  - We can't even tell the sign!
  - It looks small
  - We are confident that it's small
- The first is most common

Small effects

- Evidence that an effect is small:
    - We can’t even tell the sign!
    - It looks small
    - We are confident that it’s small
- The first is most common
    - – and the worst

Small effects

- Evidence that an effect is small:
    - We can't even tell the sign!
    - It looks small
    - We are confident that it's small
- The first is most common
    - – and the worst

Flu masks

⚠WARNING
This respirator helps protect
against certain particles.
Misuse may result in sickness
or death. For proper use, see
supervisor or box or call 3M
1-800-247-3941 NIOSH
3M 8210
N95
LOT# 49082

Flu mask example

- People who work in respiratory clinics sometimes have to wear bulky, uncomfortable, expensive masks

Flu mask example

- People who work in respiratory clinics sometimes have to wear bulky, uncomfortable, expensive masks
- They would like to switch to simpler masks, if those will do the job

Flu mask example

- People who work in respiratory clinics sometimes have to wear bulky, uncomfortable, expensive masks
- They would like to switch to simpler masks, if those will do the job
- How can this be tested statistically? We don't want the masks to be "different".

Flu mask example

- People who work in respiratory clinics sometimes have to wear bulky, uncomfortable, expensive masks
- They would like to switch to simpler masks, if those will do the job
- How can this be tested statistically? We don't want the masks to be "different".
  - Use a confidence interval

Flu mask example

- People who work in respiratory clinics sometimes have to wear bulky, uncomfortable, expensive masks
- They would like to switch to simpler masks, if those will do the job
- How can this be tested statistically? We don't want the masks to be "different".
  - Use a confidence interval
  - Decide how big a level is acceptable, and construct a P value for the hypothesis that this level is excluded!

Flu mask example

- People who work in respiratory clinics sometimes have to wear bulky, uncomfortable, expensive masks
- They would like to switch to simpler masks, if those will do the job
- How can this be tested statistically? We don't want the masks to be "different".
  - Use a confidence interval
  - Decide how big a level is acceptable, and construct a P value for the hypothesis that this level is excluded!

Study results

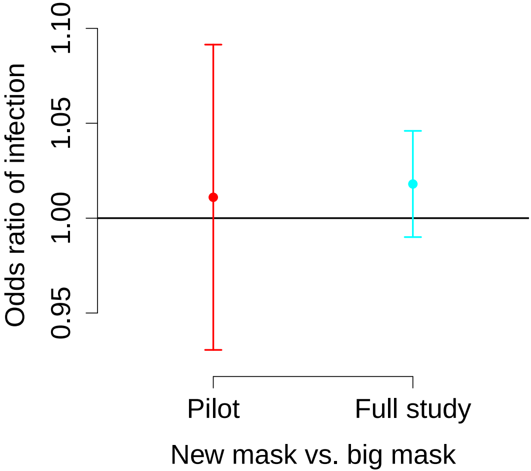

Non-inferiority trial

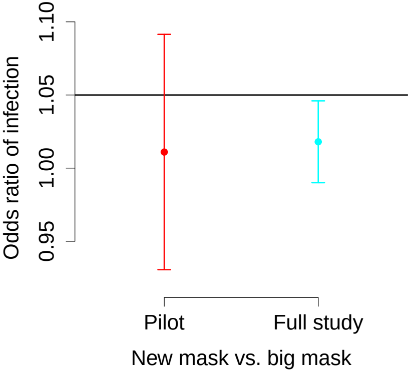

- Is the new mask “good enough”?

Non-inferiority trial

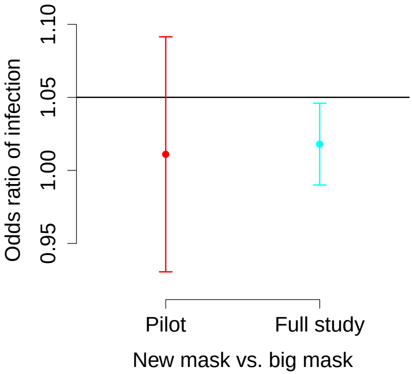

- Is the new mask “good enough”?
- What’s our standard for that?

Non-inferiority trial

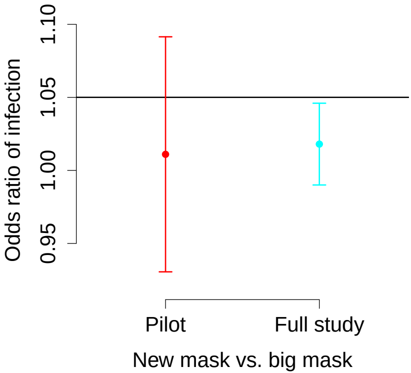

- Is the new mask “good enough”?
- What’s our standard for that?

Non-inferiority trial

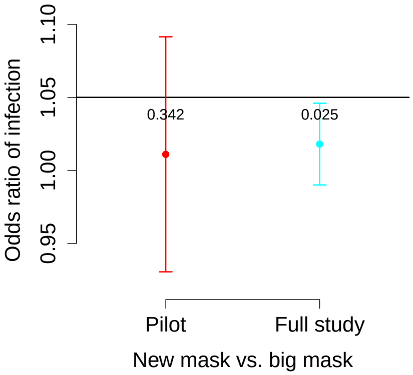

- We can even attach a P value by basing it on the “right” statistic.

Non-inferiority trial

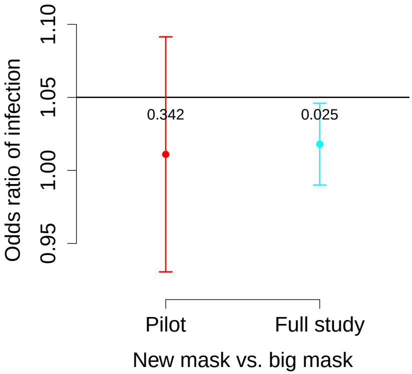

- We can even attach a P value by basing it on the “right” statistic.
- The right statistic is the thing whose sign we want to know:

Non-inferiority trial

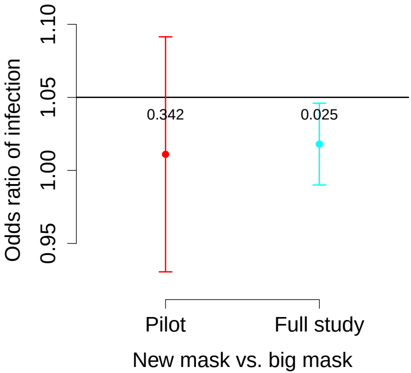

- We can even attach a P value by basing it on the “right” statistic.
- The right statistic is the thing whose sign we want to know:
    - The difference between the observed effect and the standard we chose

Non-inferiority trial

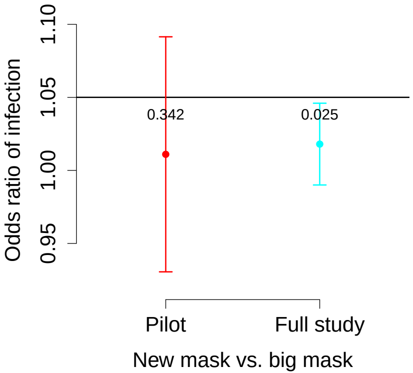

- We can even attach a P value by basing it on the "right" statistic.
- The right statistic is the thing whose sign we want to know:
  - The difference between the observed effect and the standard we chose

Making decisions

Differences

- **Never** say: A is significant and B isn't, so $A > B$

Differences

- **Never** say: A is significant and B isn't, so $A > B$
- **Instead:** Construct a statistic for the hypothesis $A > B$

Differences

- **Never** say: A is significant and B isn't, so $A > B$
- **Instead:** Construct a statistic for the hypothesis $A > B$
  - Reparameterize, study interactions

Differences

- **Never** say: A is significant and B isn't, so $A > B$
- **Instead:** Construct a statistic for the hypothesis $A > B$
  - Reparameterize, study interactions
- The Difference Between "Significant" and "Not Significant" is not Itself Statistically Significant

Differences

- **Never** say: A is significant and B isn't, so $A > B$
- **Instead:** Construct a statistic for the hypothesis $A > B$
  - Reparameterize, study interactions
- The Difference Between "Significant" and "Not Significant" is not Itself Statistically Significant
  - *Gelman and Stern*

Differences

- **Never** say: A is significant and B isn't, so $A > B$
- **Instead:** Construct a statistic for the hypothesis $A > B$
  - Reparameterize, study interactions
- The Difference Between "Significant" and "Not Significant" is not Itself Statistically Significant
  - *Gelman and Stern*

Model simplification

- It's not OK to use high P values as a standard for simplifying models

Model simplification

- It's not OK to use high P values as a standard for simplifying models
- So how do we simplify?

Model simplification

- It's not OK to use high P values as a standard for simplifying models
- So how do we simplify?
  - For prediction: information criteria

Model simplification

- It's not OK to use high P values as a standard for simplifying models
- So how do we simplify?
  - For prediction: information criteria
  - For inference: ???

Model simplification

- It's not OK to use high P values as a standard for simplifying models
- So how do we simplify?
  - For prediction: information criteria
  - For inference: ???
    - A priori approaches (including Bayesian priors)

Model simplification

- It's not OK to use high P values as a standard for simplifying models
- So how do we simplify?
  - For prediction: information criteria
  - For inference: ???
    - A priori approaches (including Bayesian priors)
    - Experiments

Model simplification

- It's not OK to use high P values as a standard for simplifying models
- So how do we simplify?
  - For prediction: information criteria
  - For inference: ???
    - A priori approaches (including Bayesian priors)
    - Experiments

Big data

- P values are rarely good for filtering

Big data

- P values are rarely good for filtering
  - We usually want to know what's big or biologically important

Big data

- P values are rarely good for filtering
  - We usually want to know what's big or biologically important
  - Not what we've seen clearly

Big data

- P values are rarely good for filtering
  - We usually want to know what's big or biologically important
  - Not what we've seen clearly
- Beware of approaches that calculate many P values in parallel

Big data

- P values are rarely good for filtering
  - We usually want to know what's big or biologically important
  - Not what we've seen clearly
- Beware of approaches that calculate many P values in parallel
- This is the cult of the P value (all statistics must be based on P values)

Big data

- P values are rarely good for filtering
    - We usually want to know what's big or biologically important
    - Not what we've seen clearly
- Beware of approaches that calculate many P values in parallel
- This is the cult of the P value (all statistics must be based on P values)

Language

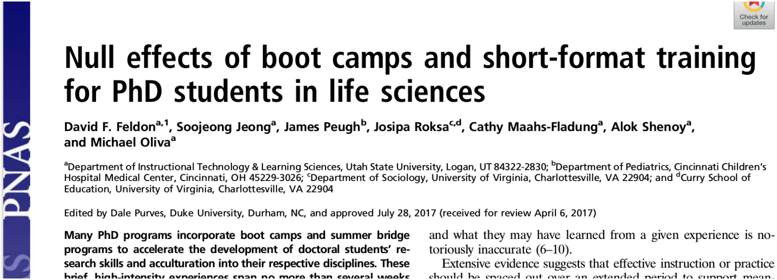

Null effects of boot camps and short-format training for PhD students in life sciences

David F. Feldon[a,1], Soojeong Jeong[a], James Peugh[b], Josipa Roksa[c,d], Cathy Maahs-Fladung[a], Alok Shenoy[a], and Michael Oliva[a]

[a]Department of Instructional Technology & Learning Sciences, Utah State University, Logan, UT 84322-2830; [b]Department of Pediatrics, Cincinnati Children's Hospital Medical Center, Cincinnati, OH 45229-3026; [c]Department of Sociology, University of Virginia, Charlottesville, VA 22904; and [d]Curry School of Education, University of Virginia, Charlottesville, VA 22904

Edited by Dale Purves, Duke University, Durham, NC, and approved July 28, 2017 (received for review April 6, 2017)

Many PhD programs incorporate boot camps and summer bridge programs to accelerate the development of doctoral students' research skills and acculturation into their respective disciplines. These brief, high-intensity experiences span no more than several weeks

and what they may have learned from a given experience is notoriously inaccurate (6–10).

Extensive evidence suggests that effective instruction or practice should be spaced out over an extended period to support mean-

Bad language

- Null effects of boot camps

Bad language

- Null effects of boot camps
  - *

Bad language

- Null effects of boot camps
    - * Wrong!

Bad language

- Null effects of boot camps
  - * Wrong!

- Lung capacity in deer-mouse populations is not correlated with elevation

Bad language

- Null effects of boot camps
  - * Wrong!
- Lung capacity in deer-mouse populations is not correlated with elevation
  - *

Bad language

- Null effects of boot camps
  - * Wrong!
- Lung capacity in deer-mouse populations is not correlated with elevation
  - * Yes, it is!

Bad language

- Null effects of boot camps
  - * Wrong!
- Lung capacity in deer-mouse populations is not correlated with elevation
  - * Yes, it is!
- As expected, the placebo group did not differ significantly from the control group

Bad language

- Null effects of boot camps
  - * Wrong!
- Lung capacity in deer-mouse populations is not correlated with elevation
  - * Yes, it is!
- As expected, the placebo group did not differ significantly from the control group
  - *

Bad language

- Null effects of boot camps
  - * Wrong!
- Lung capacity in deer-mouse populations is not correlated with elevation
  - * Yes, it is!
- As expected, the placebo group did not differ significantly from the control group
  - * Why would that be good?

Bad language

- Null effects of boot camps
  - * Wrong!
- Lung capacity in deer-mouse populations is not correlated with elevation
  - * Yes, it is!
- As expected, the placebo group did not differ significantly from the control group
  - * Why would that be good?
- B and B showed that there is no statistically significant difference in sexual risk behaviour between men with and without clinic access in Zambia

Bad language

- Null effects of boot camps
  - * Wrong!
- Lung capacity in deer-mouse populations is not correlated with elevation
  - * Yes, it is!
- As expected, the placebo group did not differ significantly from the control group
  - * Why would that be good?
- B and B showed that there is no statistically significant difference in sexual risk behaviour between men with and without clinic access in Zambia
  - *

Bad language

- Null effects of boot camps
  - * Wrong!
- Lung capacity in deer-mouse populations is not correlated with elevation
  - * Yes, it is!
- As expected, the placebo group did not differ significantly from the control group
  - * Why would that be good?
- B and B showed that there is no statistically significant difference in sexual risk behaviour between men with and without clinic access in Zambia
  - * No, they didn't

Bad language

- Null effects of boot camps
  - * Wrong!
- Lung capacity in deer-mouse populations is not correlated with elevation
  - * Yes, it is!
- As expected, the placebo group did not differ significantly from the control group
  - * Why would that be good?
- B and B showed that there is no statistically significant difference in sexual risk behaviour between men with and without clinic access in Zambia
  - * No, they didn't
  - *

Bad language

- Null effects of boot camps
  - * Wrong!
- Lung capacity in deer-mouse populations is not correlated with elevation
  - * Yes, it is!
- As expected, the placebo group did not differ significantly from the control group
  - * Why would that be good?
- B and B showed that there is no statistically significant difference in sexual risk behaviour between men with and without clinic access in Zambia
  - * No, they didn't
  - * Statistical significance is a property of the *study* (and the sample population), never of real groups (or the idealized population)

Bad language

- Null effects of boot camps
  - * Wrong!
- Lung capacity in deer-mouse populations is not correlated with elevation
  - * Yes, it is!
- As expected, the placebo group did not differ significantly from the control group
  - * Why would that be good?
- B and B showed that there is no statistically significant difference in sexual risk behaviour between men with and without clinic access in Zambia
  - * No, they didn't
  - * Statistical significance is a property of the *study* (and the sample population), never of real groups (or the idealized population)

Confusion

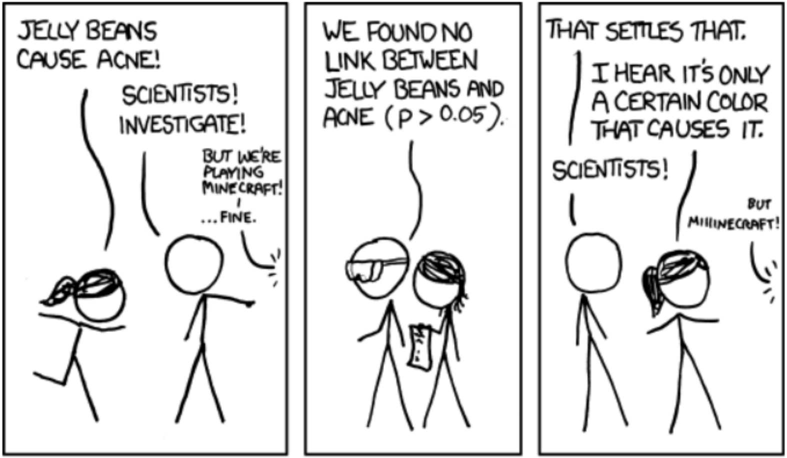
JELLY BEANS CAUSE ACNE!
SCIENTISTS! INVESTIGATE!
BUT WE'RE PLAYING MINECRAFT!
...FINE.
WE FOUND NO LINK BETWEEN JELLY BEANS AND ACNE (P > 0.05).
THAT SETTLES THAT.
I HEAR IT'S ONLY A CERTAIN COLOR THAT CAUSES IT.
SCIENTISTS!
BUT MIIIINECRAFT!

Improving language

- **Wrong:** This treatment does not have a statistically significant effect

Improving language

- **Wrong:** This treatment does not have a statistically significant effect
- **Standard:** We found that this treatment has no statistically significant effect

Improving language

- **Wrong:** This treatment does not have a statistically significant effect
- **Standard:** We found that this treatment has no statistically significant effect
- **Better:** We did not find a statistically significant effect of this treatment

Improving language

- **Wrong:** This treatment does not have a statistically significant effect
- **Standard:** We found that this treatment has no statistically significant effect
- **Better:** We did not find a statistically significant effect of this treatment
- **Best** ??

Improving language

- **Wrong:** This treatment does not have a statistically significant effect
- **Standard:** We found that this treatment has no statistically significant effect
- **Better:** We did not find a statistically significant effect of this treatment
- **Best** ??

Is statistical “significance” a thing?

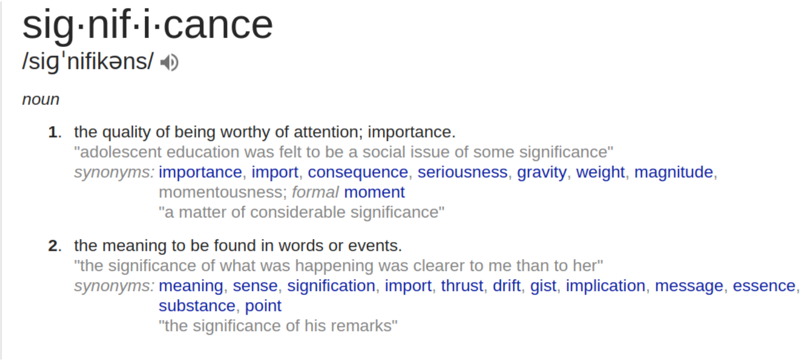

sig·nif·i·cance

/sigˈnifikəns/

noun

1. the quality of being worthy of attention; importance.
   "adolescent education was felt to be a social issue of some significance"
   synonyms: importance, import, consequence, seriousness, gravity, weight, magnitude, momentousness; *formal* moment
   "a matter of considerable significance"

2. the meaning to be found in words or events.
   "the significance of what was happening was clearer to me than to her"
   synonyms: meaning, sense, signification, import, thrust, drift, gist, implication, message, essence, substance, point
   "the significance of his remarks"

- *

Is statistical “significance” a thing?

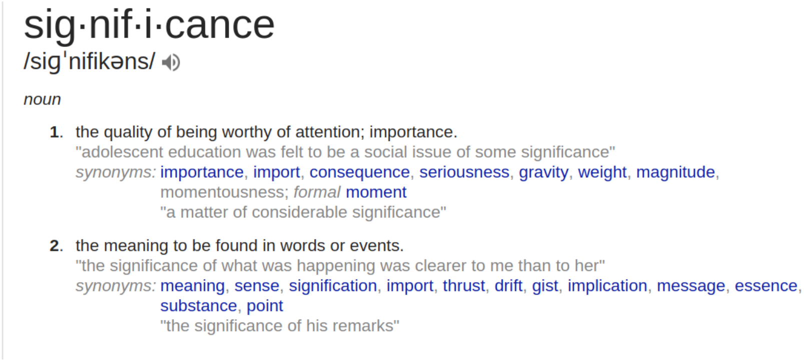

sig·nif·i·cance

/sigˈnifikəns/ 🔊

noun

1. the quality of being worthy of attention; importance.
   "adolescent education was felt to be a social issue of some significance"
   synonyms: importance, import, consequence, seriousness, gravity, weight, magnitude, momentousness; *formal* moment
   "a matter of considerable significance"

2. the meaning to be found in words or events.
   "the significance of what was happening was clearer to me than to her"
   synonyms: meaning, sense, signification, import, thrust, drift, gist, implication, message, essence, substance, point
   "the significance of his remarks"

- * It may be a thing

Is statistical “significance” a thing?

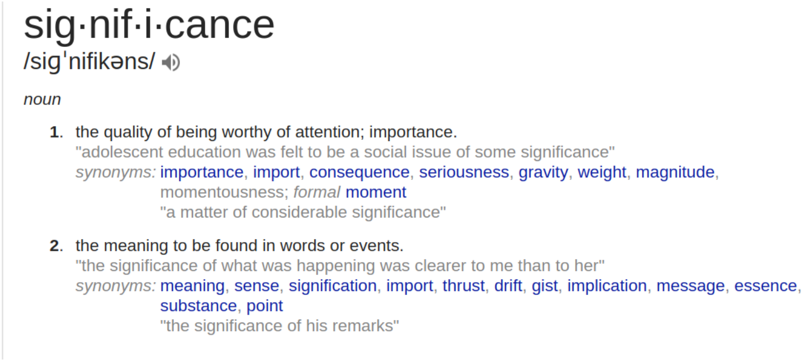

sig·nif·i·cance

/sigˈnifikəns/

noun

1. the quality of being worthy of attention; importance.
   "adolescent education was felt to be a social issue of some significance"
   synonyms: importance, import, consequence, seriousness, gravity, weight, magnitude, momentousness; *formal* moment
   "a matter of considerable significance"

2. the meaning to be found in words or events.
   "the significance of what was happening was clearer to me than to her"
   synonyms: meaning, sense, signification, import, thrust, drift, gist, implication, message, essence, substance, point
   "the significance of his remarks"

- * It may be a thing
- *

Is statistical "significance" a thing?

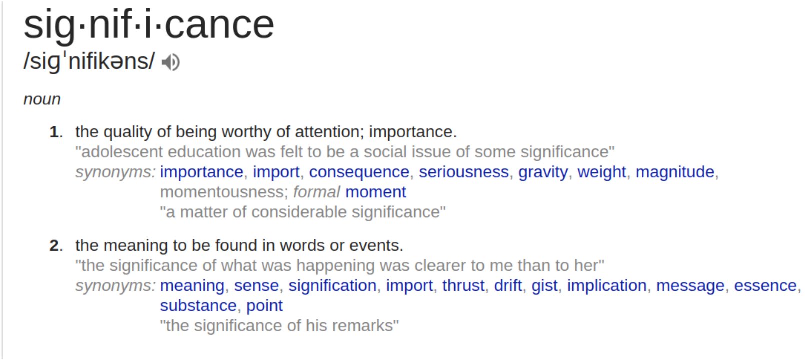

sig·nif·i·cance

/sigˈnifikəns/

noun

1. the quality of being worthy of attention; importance.
   "adolescent education was felt to be a social issue of some significance"
   synonyms: importance, import, consequence, seriousness, gravity, weight, magnitude, momentousness; *formal* moment
   "a matter of considerable significance"

2. the meaning to be found in words or events.
   "the significance of what was happening was clearer to me than to her"
   synonyms: meaning, sense, signification, import, thrust, drift, gist, implication, message, essence, substance, point
   "the significance of his remarks"

- * It may be a thing
- * But it's not much to do with the normal meaning of significance

Is statistical "significance" a thing?

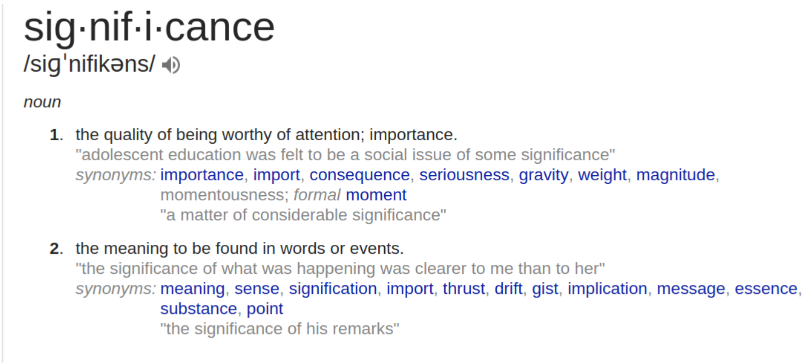

sig·nif·i·cance

/sigˈnifikəns/

noun

1. the quality of being worthy of attention; importance.
   "adolescent education was felt to be a social issue of some significance"
   synonyms: importance, import, consequence, seriousness, gravity, weight, magnitude, momentousness; *formal* moment
   "a matter of considerable significance"
2. the meaning to be found in words or events.
   "the significance of what was happening was clearer to me than to her"
   synonyms: meaning, sense, signification, import, thrust, drift, gist, implication, message, essence, substance, point
   "the significance of his remarks"

- * It may be a thing
- * But it's not much to do with the normal meaning of significance

Fish hormones

- Male fish subject to polluted water have more female hormones than controls

Fish hormones

- Male fish subject to polluted water have more female hormones than controls
  - $P < 0.05$

Fish hormones

- Male fish subject to polluted water have more female hormones than controls
  - $P < 0.05$
  - A significant effect!

Fish hormones

- Male fish subject to polluted water have more female hormones than controls
  - $P < 0.05$
  - A significant effect!
- Is it a significant amount of hormone? How much hormone is it?

Fish hormones

- Male fish subject to polluted water have more female hormones than controls
  - $P<0.05$
  - A significant effect!
- Is it a significant amount of hormone? How much hormone is it?

What do P values measure?

- *

What do P values measure?

- * Clarity!

What do P values measure?

- * Clarity!

- *

What do P values measure?

- * Clarity!
- * We should call it that

What do P values measure?

- * Clarity!
- * We should call it that

Another way to talk

- As expected, the sign of the difference between the placebo group and controls was unclear

Another way to talk

- As expected, the sign of the difference between the placebo group and controls was unclear
- Unclear effects of boot camps

Another way to talk

- As expected, the sign of the difference between the placebo group and controls was unclear
- Unclear effects of boot camps
- The direction of correlation between lung capacity and elevation in deer-mouse populations is unclear

Another way to talk

- As expected, the sign of the difference between the placebo group and controls was unclear
- Unclear effects of boot camps
- The direction of correlation between lung capacity and elevation in deer-mouse populations is unclear
- B and B showed that there is an unclear difference in sexual risk behaviour between men with and without clinic access in Zambia

Another way to talk

- As expected, the sign of the difference between the placebo group and controls was unclear
- Unclear effects of boot camps
- The direction of correlation between lung capacity and elevation in deer-mouse populations is unclear
- B and B showed that there is an unclear difference in sexual risk behaviour between men with and without clinic access in Zambia

Improving language

- **Wrong:** This treatment does not have a statistically significant effect

Improving language

- **Wrong:** This treatment does not have a statistically significant effect
- **Standard:** We found that this treatment has no statistically significant effect

Improving language

- **Wrong:** This treatment does not have a statistically significant effect
- **Standard:** We found that this treatment has no statistically significant effect
- **Better:** We did not find a statistically significant effect of this treatment

Improving language

- **Wrong:** This treatment does not have a statistically significant effect
- **Standard:** We found that this treatment has no statistically significant effect
- **Better:** We did not find a statistically significant effect of this treatment
- **New:** We did not *see* a statistically *clear* effect of this treatment

Improving language

- **Wrong:** This treatment does not have a statistically significant effect
- **Standard:** We found that this treatment has no statistically significant effect
- **Better:** We did not find a statistically significant effect of this treatment
- **New:** We did not *see* a statistically *clear* effect of this treatment
  - The effect of this treatment was not statistically *clear* in this study

Improving language

- **Wrong:** This treatment does not have a statistically significant effect
- **Standard:** We found that this treatment has no statistically significant effect
- **Better:** We did not find a statistically significant effect of this treatment
- **New:** We did not *see* a statistically *clear* effect of this treatment
  - The effect of this treatment was not statistically *clear* in this study

Is it possible?

- It's hard to get people to change language

Is it possible?

- It's hard to get people to change language
- But you can probably change your language (if you keep the P values)

Is it possible?

- It's hard to get people to change language
- But you can probably change your language (if you keep the P values)
  - We found a statistically clear increase (P=0.02) in blood iron in the vitamin-supplement group

Is it possible?

- It's hard to get people to change language
- But you can probably change your language (if you keep the P values)
  - We found a statistically clear increase (P=0.02) in blood iron in the vitamin-supplement group
  - The direction of association between lung capacity and elevation was not statistically clear (P=0.43)

Is it possible?

- It's hard to get people to change language
- But you can probably change your language (if you keep the P values)
  - We found a statistically clear increase (P=0.02) in blood iron in the vitamin-supplement group
  - The direction of association between lung capacity and elevation was not statistically clear (P=0.43)
  - B and B did not see a statistically clear difference in sexual risk behaviour between men with and without clinic access in Zambia (P=0.1)

Is it possible?

- It's hard to get people to change language
- But you can probably change your language (if you keep the P values)
    - We found a statistically clear increase (P=0.02) in blood iron in the vitamin-supplement group
    - The direction of association between lung capacity and elevation was not statistically clear (P=0.43)
    - B and B did not see a statistically clear difference in sexual risk behaviour between men with and without clinic access in Zambia (P=0.1)
- Confidence intervals are still better, when possible

Is it possible?

- It's hard to get people to change language
- But you can probably change your language (if you keep the P values)
  - We found a statistically clear increase (P=0.02) in blood iron in the vitamin-supplement group
  - The direction of association between lung capacity and elevation was not statistically clear (P=0.43)
  - B and B did not see a statistically clear difference in sexual risk behaviour between men with and without clinic access in Zambia (P=0.1)
- Confidence intervals are still better, when possible

Statistical philosophy

Advice for scientists

- Statistics are not a magic machine that gives you the right answer

Statistical philosophy

Advice for scientists

- Statistics are not a magic machine that gives you the right answer
- If you are to be a serious scientist in a noisy world, you should have your own philosophy of statistics

Statistical philosophy

Advice for scientists

- Statistics are not a magic machine that gives you the right answer
- If you are to be a serious scientist in a noisy world, you should have your own philosophy of statistics
  - Be pragmatic: your goal is to do science, not get caught by theoretical considerations

Statistical philosophy

Advice for scientists

- Statistics are not a magic machine that gives you the right answer
- If you are to be a serious scientist in a noisy world, you should have your own philosophy of statistics
  - Be pragmatic: your goal is to do science, not get caught by theoretical considerations
  - Be honest: it's harder than it sounds.

Statistical philosophy

Advice for scientists

- Statistics are not a magic machine that gives you the right answer
- If you are to be a serious scientist in a noisy world, you should have your own philosophy of statistics
    - Be pragmatic: your goal is to do science, not get caught by theoretical considerations
    - Be honest: it's harder than it sounds.

Honesty

- You can always keep analyzing until you find a “significant” (or “clear”) result

Honesty

- You can always keep analyzing until you find a “significant” (or “clear”) result
  - If you do this you will make a lot of mistakes

Honesty

- You can always keep analyzing until you find a “significant” (or “clear”) result
  - If you do this you will make a lot of mistakes
- You may also keep analyzing until you find a result that you already “know” is true.

Honesty

- You can always keep analyzing until you find a “significant” (or “clear”) result
  - If you do this you will make a lot of mistakes
- You may also keep analyzing until you find a result that you already “know” is true.
  - This is confirmation bias; you’re probably right, but your project is not advancing science

Honesty

- You can always keep analyzing until you find a “significant” (or “clear”) result
  - If you do this you will make a lot of mistakes
- You may also keep analyzing until you find a result that you already “know” is true.
  - This is confirmation bias; you’re probably right, but your project is not advancing science
- Good practice

Honesty

- You can always keep analyzing until you find a “significant” (or “clear”) result
  - If you do this you will make a lot of mistakes
- You may also keep analyzing until you find a result that you already “know” is true.
  - This is confirmation bias; you’re probably right, but your project is not advancing science
- Good practice
  - Keep a data-analysis journal

Honesty

- You can always keep analyzing until you find a “significant” (or “clear”) result
  - If you do this you will make a lot of mistakes
- You may also keep analyzing until you find a result that you already “know” is true.
  - This is confirmation bias; you’re probably right, but your project is not advancing science
- Good practice
  - Keep a data-analysis journal
  - Start *before* you look at the data

Honesty

- You can always keep analyzing until you find a “significant” (or “clear”) result
    - If you do this you will make a lot of mistakes
- You may also keep analyzing until you find a result that you already “know” is true.
    - This is confirmation bias; you’re probably right, but your project is not advancing science
- Good practice
    - Keep a data-analysis journal
    - Start *before* you look at the data

Summary

- P values are over-rated

Summary

- P values are over-rated
- High P values should not be used as evidence for anything ever.

Summary

- P values are over-rated
- High P values should not be used as evidence for anything ever.
    - They can provide indirect evidence. Wonderful. Find the direct evidence and use that instead.

Summary

- P values are over-rated
- High P values should not be used as evidence for anything ever.
  - They can provide indirect evidence. Wonderful. Find the direct evidence and use that instead.
- Use effect sizes and confidence intervals when you can

Summary

- P values are over-rated
- High P values should not be used as evidence for anything ever.
    - They can provide indirect evidence. Wonderful. Find the direct evidence and use that instead.
- Use effect sizes and confidence intervals when you can
- Otherwise, find ways to make low P values do the work

Summary

- P values are over-rated
- High P values should not be used as evidence for anything ever.
  - They can provide indirect evidence. Wonderful. Find the direct evidence and use that instead.
- Use effect sizes and confidence intervals when you can
- Otherwise, find ways to make low P values do the work
  - Non-inferiority tests, interactions

Summary

- P values are over-rated
- High P values should not be used as evidence for anything ever.
  - They can provide indirect evidence. Wonderful. Find the direct evidence and use that instead.
- Use effect sizes and confidence intervals when you can
- Otherwise, find ways to make low P values do the work
  - Non-inferiority tests, interactions
  - Don't rely on unclear information

Summary

- P values are over-rated
- High P values should not be used as evidence for anything ever.
  - They can provide indirect evidence. Wonderful. Find the direct evidence and use that instead.
- Use effect sizes and confidence intervals when you can
- Otherwise, find ways to make low P values do the work
  - Non-inferiority tests, interactions
  - Don't rely on unclear information

Language

- Language is important and feeds misunderstanding

Language

- Language is important and feeds misunderstanding
- Even if you are not misled, others will be

Language

- Language is important and feeds misunderstanding
- Even if you are not misled, others will be
- Use language clearly:

Language

- Language is important and feeds misunderstanding
- Even if you are not misled, others will be
- Use language clearly:
  - We found no difference

Language

- Language is important and feeds misunderstanding
- Even if you are not misled, others will be
- Use language clearly:
  - We found no difference
  - $\Longrightarrow$ We did not see a clear difference

Language

- Language is important and feeds misunderstanding
- Even if you are not misled, others will be
- Use language clearly:
  - We found no difference
  - $\Longrightarrow$ We did not see a clear difference
- Consider abandoning the language of statistical “significance”

Language

- Language is important and feeds misunderstanding
- Even if you are not misled, others will be
- Use language clearly:
  - We found no difference
  - $\Longrightarrow$ We did not see a clear difference
- Consider abandoning the language of statistical “significance”
- *Definitely* abandon the language of statistical “equivalence”

Language

- Language is important and feeds misunderstanding
- Even if you are not misled, others will be
- Use language clearly:
  - We found no difference
  - $\Longrightarrow$ We did not see a clear difference
- Consider abandoning the language of statistical “significance”
- *Definitely* abandon the language of statistical “equivalence”